

CSCI-1200 Data Structures — Fall 2017

Lab 12 — Order Notation & Performance

In this lab we will carry out a series of tests on the five fundamental data structures in the Standard Template Library we have studied this semester: `vector`, `list`, binary search tree (`set` / `map`), `priority_queue`, and hash table (`unordered_set` / `unordered_map`). We will hypothesize and evaluate the relative performance of these data structures and solidify our understanding of order notation.

For each data structure you will predict and measure the runtime for three different moderately compute-intensive operations: *sorting*, *removing duplicates* (without changing the overall order), and *finding the mode* (most frequently occurring element). You will perform these tests on STL `string` objects that can be read from a file or constructed from random sequences of `chars`.

Checkpoint 1

Before doing any implementation or testing, create a simple table with 5 data structures on one axis and 3 operations on the other axis. In each cell write down the order notation for completing the operation using that datatype. If it is not feasible/sensible to use a particular data structure to complete the specified operation put an X in the box. If two data structures have the same order notation for one operation, predict which one will have a faster running time for large data. We combine `set` & `map` (and `unordered_set` & `unordered_map`) in this analysis, but be sure to specify which datatype of the two makes the most sense for each operation.

To complete this checkpoint: Present your analysis to one of the TAs and be prepared to explain your answers and an outline/pseudocode of the implementation as necessary.

Checkpoint 2

Next, let's dive into the provided code. We have implemented the 3 operations for the `vector` datatype.

```
http://www.cs.rpi.edu/academics/courses/fall17/csci1200/labs/12\_order\_notation/performance.cpp  
http://www.cs.rpi.edu/academics/courses/fall17/csci1200/labs/12\_order\_notation/test\_input.txt  
http://www.cs.rpi.edu/academics/courses/fall17/csci1200/labs/12\_order\_notation/test\_sort.txt  
http://www.cs.rpi.edu/academics/courses/fall17/csci1200/labs/12\_order\_notation/test\_mode.txt  
http://www.cs.rpi.edu/academics/courses/fall17/csci1200/labs/12\_order\_notation/test\_remove\_dups.txt
```

The data structure, operation, source of the input, size & length of input (for randomly generated input), and output file are specified on the command line. Here are a few sample command lines:

```
a.exe vector mode in.txt out.txt  
a.exe vector sort random 10000 5 out.txt  
a.exe vector remove_dups random 10000 2 out.txt
```

The first example reads the file named `in.txt`, uses a vector to find the most frequently occurring value (implemented by first sorting the data), and then outputs that string (the mode) to a file named `out.txt`. The second example will generate 10,000 random strings of length 5, use a vector to sort them, and then output the result to the file named `out.txt`. Similarly, the command line option `remove_dups` will remove the duplicate elements from the input (without otherwise changing the order).

Compile and test the provided code on a variety of tests of different sizes for the `vector` datatype for each of the 3 operations. The provided code uses the `clock()` function to measure the processing time of the computation. The resolution accuracy of the timing mechanism is system and hardware dependent and may be in seconds, milliseconds, or something else. If the resolution on your system is coarse, you must base your analysis on the measurements from sufficiently large datasets. The program reports the time to load, process, and save the data. Record the results in a table like this:

Sorting random 5 letter strings using STL vector

# of strings	load time (sec)	operation time (sec)	output time (sec)
10000	0.023	0.031	0.089
20000	0.043	0.067	0.172
50000	0.115	0.180	0.445
100000	0.226	0.402	0.918

As the dataset grows, does your predicted order notation match the raw performance numbers? For example, in the table above, the time to both load and output is linear, in the # of strings, that is $O(n)$.

$$\begin{aligned} \text{load time}(n) &= k_{\text{load}} * n \\ \text{output time}(n) &= k_{\text{output}} * n \end{aligned}$$

We can estimate the coefficients from the collected numbers: $k_{\text{load}} \approx 2.2 \times 10^{-6}$ and $k_{\text{output}} \approx 9.0 \times 10^{-6}$. The running time for sorting with the STL **vector** sorting algorithm is $O(n \log_2 n)$:

$$\text{operation time}(n) = k_{\text{operation}} * n \log_2 n$$

with coefficient $k_{\text{operation}} \approx 2.3 \times 10^{-7}$. Of course these constants will be different on different operating systems, different compilers, and different hardware! These constants will allow us to compare data structures / algorithms with the same order notation.

To complete this checkpoint: Show the results of your testing and analysis to one of the TAs. The data should be neat & well organized in order to receive this checkpoint.